

Applied Cryptography

Spring Semester 2023

Lectures 22, 23, 24 and 25

Kenny Paterson (@kennyog)

Applied Cryptography Group

<https://appliedcrypto.ethz.ch/>

Overview of lectures

- Introducing Public Key Encryption
- The KEM/DEM paradigm
- RSA Encryption
- Diffie-Hellman Key Exchange and ElGamal Encryption
- Appendix: Basic Number Theory

Introducing Public Key Encryption

Introducing Public Key Encryption

- So far, we have considered symmetric encryption schemes.
- Participants need to somehow share key K in order to encrypt/decrypt.
- In Public Key Encryption (PKE), we use different keys for **encryption** and **decryption**.
 - Bob generates a **public**/**private** key pair.
 - Alice uses Bob's **public** key to **encrypt** to Bob.
 - Bob uses his **private** key to **decrypt**.
- Also called Asymmetric Encryption.

Syntax for PKE

A PKE scheme \mathcal{PKE} consists of a triple of algorithms: $\mathcal{PKE} = (\text{KGen}, \text{Enc}, \text{Dec})$.

KGen: randomised key generation, generates a key pair $(sk, pk) \in \mathcal{SK} \times \mathcal{PK}$.
(sk is the **private** key, pk the **public** key).

Enc: usually randomised, takes as input **public key** pk , plaintext $m \in \mathcal{M} \subseteq \{0, 1\}^*$ and produces output $c \in \mathcal{C} \subseteq \{0, 1\}^*$.

Dec: takes as input **private key** sk , ciphertext $c \in \{0, 1\}^*$ and produces output $m \in \mathcal{M}$, or an error message denoted \perp .

Correctness: we require that for all key pairs (sk, pk) output by KGen, and for all plaintexts m ,

$$\text{Dec}_{sk}(\text{Enc}_{pk}(m)) = m.$$

A concrete example: Textbook RSA

KGen: generates:

- random primes p, q of some bit-size $k/2$, with $N = pq$ (so bit-size of N is k);
- integers d, e such that $de = 1 \bmod (p-1)(q-1)$.

Outputs key pair (sk, pk) where $sk = d, pk = (e, N)$.

Enc: inputs public key $pk = (e, N)$, plaintext $m \in [1, N-1]$; outputs $c = m^e \bmod N$.

Dec: inputs private key $sk = d$, ciphertext c ; outputs $m = c^d \bmod N$.

Correctness: follows from property that if $de = 1 \bmod (p-1)(q-1)$ then $m^{de} = m \bmod N$ for all $m \in [1, N-1]$ (this demands a mathematical proof).

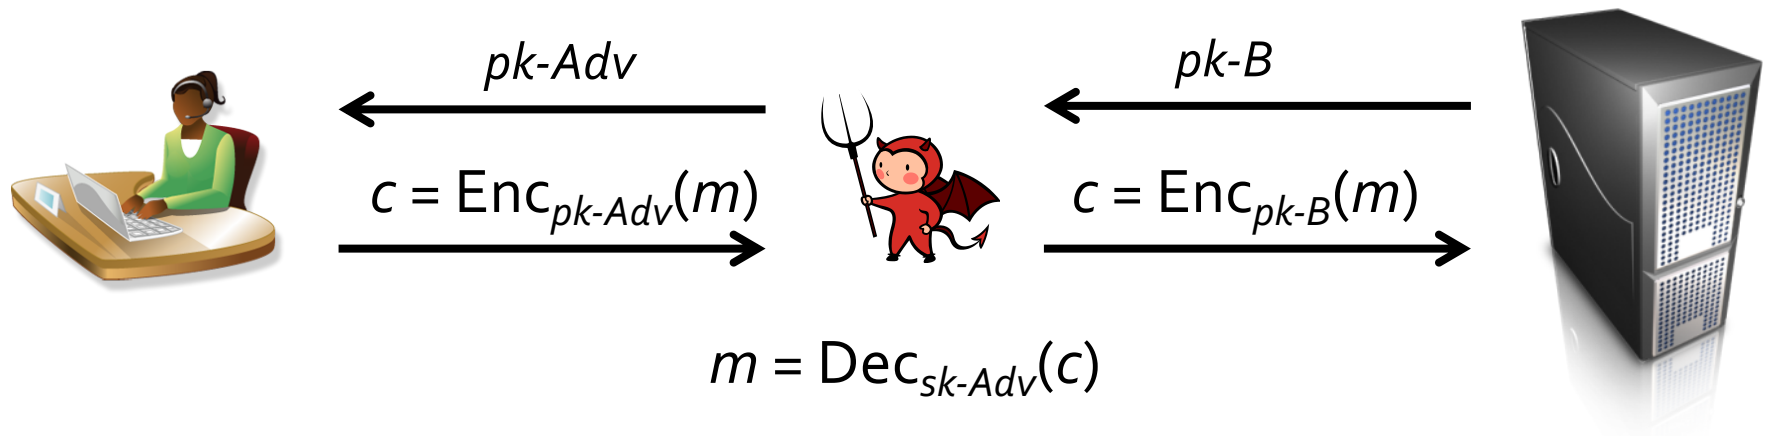
- Implementation issues: how do we generate random primes of a given bit-size? How do we generate d and e ? How do we encode messages as integers in the interval $[1, N-1]$?
- Enc is not randomised here, which already hints at trouble.
- **This scheme is not secure. It must not be used as is in practice.**

What is PKE good for?

- PKE is typically much more expensive than symmetric encryption (in speed, key-size, etc)...
- Example: for textbook RSA, encryption involves an exponentiation mod N , where N typically has 2048 bits.
- PKE is often used in applications to transport **symmetric keys which are then used to encrypt bulk data**.
- This use of PKE is often referred to as **hybrid encryption**.
 - The symmetric keys are then used in a symmetric encryption scheme (e.g. nonce-based AEAD) to encrypt the actual data.
 - We will return to this in a few slides time.
- PKE is occasionally used to directly transport short, infrequent messages (e.g. PINs from card to payment terminal in EMV system).

What is PKE good for?

- Main problem with PKE: it requires distribution of **authentic** public keys.
- How does Alice know that a public key is genuinely Bob's?



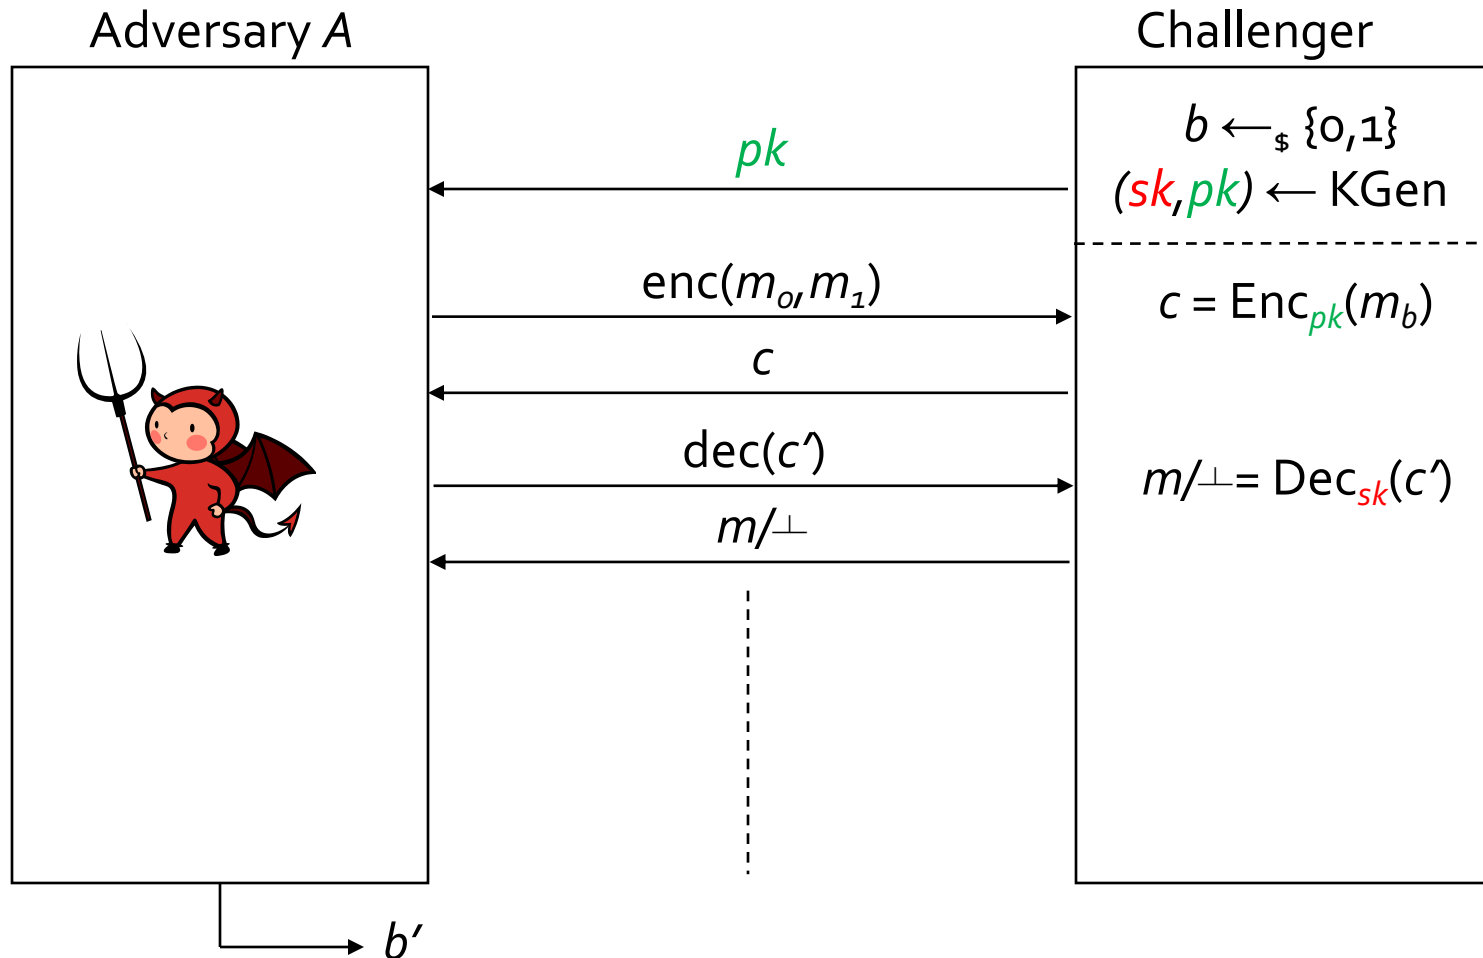
- At a high level, PKE translates the problem of symmetric key distribution into the problem of distribution of authentic public keys.
- **Question: Is that an easier problem?**

Formalising security for PKE

- As with symmetric encryption, it is not enough that the secret key should be hard to guess: we want confidentiality of messages, more precisely **indistinguishability under chosen ciphertext attacks (IND-CCA)**.
- Captures the idea that nothing about the plaintext should leak to the adversary, even when the adversary can obtain decryptions of arbitrary messages.
 - Adversary has access to an *encryption* oracle and a *decryption* oracle for key pair (sk, pk) .
 - Adversary can submit arbitrary pairs of messages (m_0, m_1) to the encryption oracle, and receives $\text{Enc}_{pk}(m_b)$.
 - Adversary can submit arbitrary* bit-strings c' to the decryption oracle and receives $\text{Dec}_{sk}(c')$.
 - Adversary has to recover the bit b .
 - Compare to definition for **IND-CCA** security of *symmetric* encryption.

*Adversary cannot submit an output from encryption oracle to decryption oracle, otherwise he can win trivially.

IND-CCA security for PKE in a picture



$$\text{Adv}_{\text{PKE}}^{\text{IND-CCA}}(A) := 2|\Pr(b=b') - 1/2|.$$

IND-CCA security for PKE

- The adversary's *advantage* in the IND-CCA security game is defined to be:

$$\text{Adv}_{\mathcal{PKE}}^{\text{IND-CCA}}(A) := 2|\Pr(b=b') - 1/2|.$$

- A PKE scheme $\mathcal{PKE} = (\text{KGen}, \text{Enc}, \text{Dec})$ is said to be **IND-CCA**-secure if the advantage is “small” for *any* adversary using “reasonable” resources.
- More precisely, we say that a scheme \mathcal{PKE} is (q_e, q_d, t, ϵ) -IND-CCA-secure if no adversary running in time t , making q_e $\text{enc}(\cdot)$ queries and q_d $\text{dec}(\cdot)$ queries has advantage more than ϵ .
- As with SE, we will provide concrete security reductions to underlying hard problems.
- One can also use complexity-theoretic notions such a poly-time, negligible, etc, after introducing a security parameter κ to dictate key sizes.
- Usually only a single $\text{enc}(\cdot)$ query is allowed; one can show equivalence to model with multiple such queries with a factor of q_e loss in security.
- IND-CPA** security: as IND-CCA but remove the decryption oracle.

Implications of the IND-CCA definition

Q: Suppose an attacker can recover sk from pk by some means for a PKE scheme. Can the scheme be IND-CCA secure?

A: No. Attacker can just decrypt $c = \text{Enc}_{pk}(m_b)$.

Q. Suppose a scheme is IND-CCA secure. Is it necessarily IND-CPA secure?

A. Yes, because removing the decryption oracle only makes the adversary less powerful. (Can easily provide a reduction proof.)

Q. Suppose a PKE scheme has a deterministic encryption algorithm. Can it be IND-CPA secure?

A. No. Make one encryption query on (m_o, m_1) to receive c ; then encrypt both m_o and m_1 directly using pk and compare ciphertexts.

Q. Is textbook RSA IND-CCA secure? IND-CPA secure?

A. It's deterministic, so it cannot achieve either notion.

Hybrid Encryption + KEM/DEM Paradigm

Key Encapsulation Mechanisms (KEMs)

A KEM \mathcal{KEM} consists of a triple of algorithms: $\mathcal{KEM} = (\text{KGen}, \text{Encap}, \text{Decap})$.

KGen: randomised key generation, generates a key pair $(sk, pk) \in \mathcal{SK} \times \mathcal{PK}$.
(sk is the **private** key, pk the **public** key).

Encap: usually randomised, takes as input **public key** pk , and produces output $(c, K) \in \mathcal{C} \times \mathcal{K}$.

Decap: takes as input **private key** sk , encapsulation $c \in \{0, 1\}^*$ and produces output $K \in \mathcal{K}$, or an error message denoted \perp .

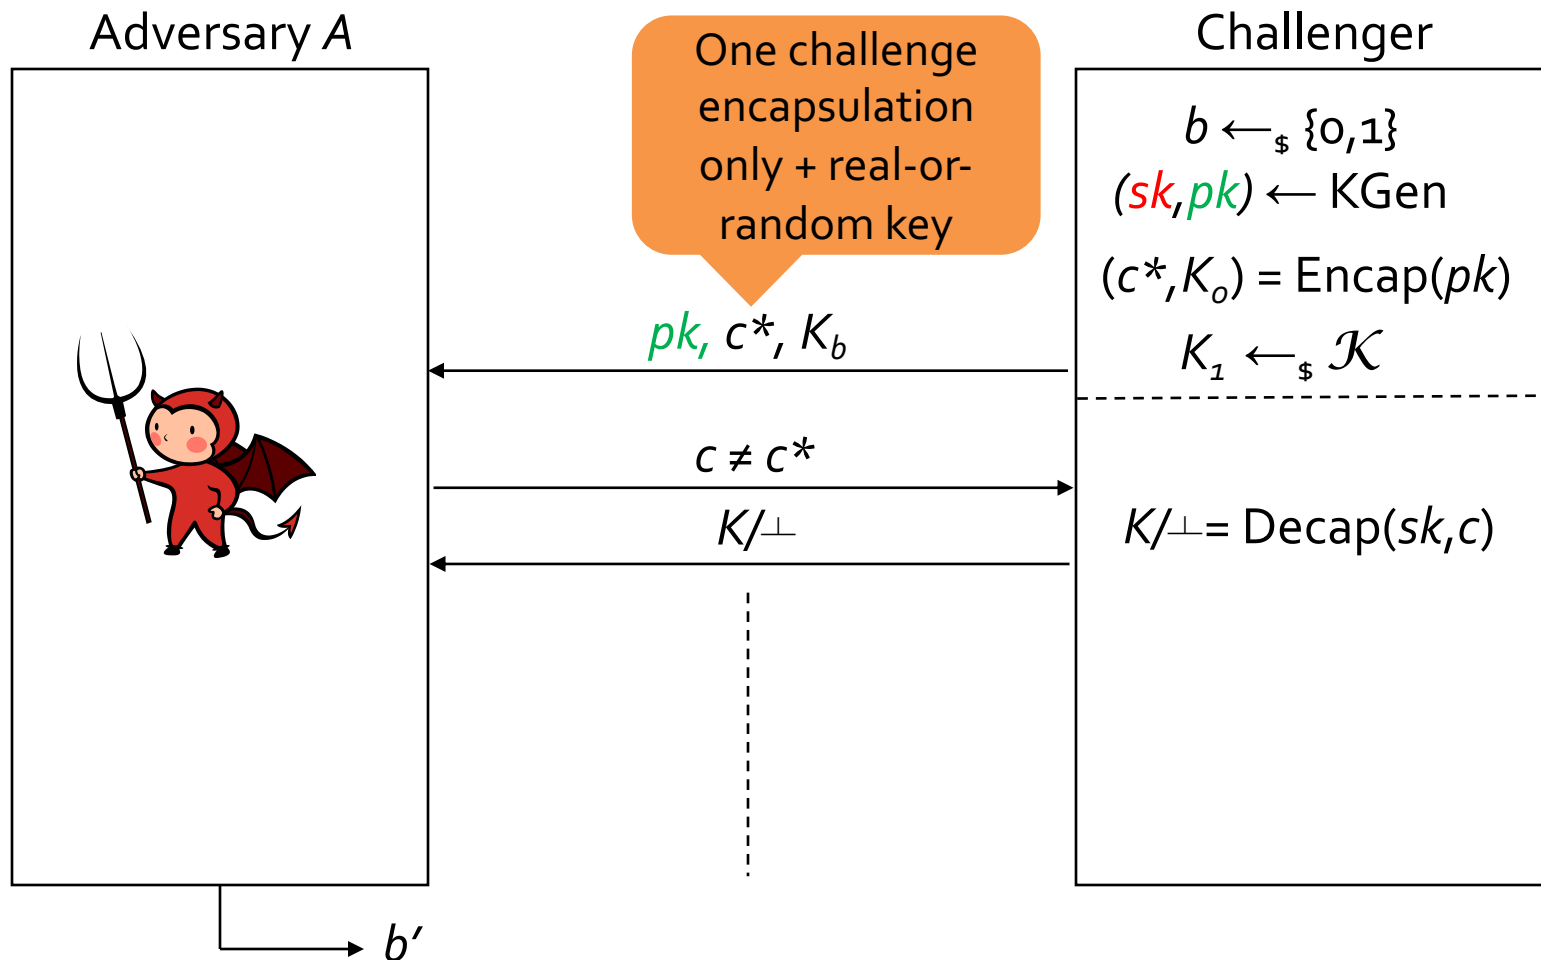
Correctness: for all key pairs (sk, pk) output by KGen,

if $(c, K) \leftarrow \text{Encap}(pk)$,

then $K \leftarrow \text{Decap}(sk, c)$.

Comparison to PKE: Encap has no message input; Encap internally generates a key K and its encapsulation c .

IND-CCA security for KEMs in a picture



The KEM/DEM Composition Paradigm

A Data Encapsulation Mechanism (DEM) is nothing other than a symmetric encryption scheme \mathcal{DEM} , with algorithms (DEM.KGen, DEM.Enc, DEM.Dec).

KEM/DEM Composition:

Let $\mathcal{KEM} = (\text{KGen}, \text{Encap}, \text{Decap})$ be a KEM and $\mathcal{DEM} = (\text{DEM.KGen}, \text{DEM.Enc}, \text{DEM.Dec})$ be a DEM such that $\mathcal{KEM}.\mathcal{K} = \mathcal{DEM}.\mathcal{K}$.

Then we build a PKE scheme $\mathcal{PKE} = (\text{PKE.KGen}, \text{PKE.Enc}, \text{PKE.Dec})$ from \mathcal{KEM} and \mathcal{DEM} as follows:

PKE.KGen: $(sk, pk) \leftarrow \text{KEM.KGen}$; **return** (sk, pk) .

PKE.Enc(pk, m):

1. $(c_o, K) \leftarrow \text{KEM.Encap}(pk)$;
2. $c_1 \leftarrow \text{DEM.Enc}(K, m)$;
3. **return** (c_o, c_1) .

PKE.Dec(sk, (c_o, c₁)):

1. $m \leftarrow \perp$;
2. $K \leftarrow \text{KEM.Decap}(sk, c_o)$;
3. If $K \neq \perp$ then $m \leftarrow \text{DEM.Dec}(K, c_1)$;
4. **return** m .

NB: message space for \mathcal{PKE} is $\mathcal{DEM}.\mathcal{M}$; ciphertext space for \mathcal{PKE} is $\mathcal{KEM}.C \times \mathcal{DEM}.C$.

The KEM/DEM Composition Paradigm

Theorem:

Suppose \mathcal{PKE} is built from \mathcal{KEM} and \mathcal{DEM} as above. If \mathcal{KEM} is IND-CCA secure and \mathcal{DEM} is IND-CCA secure, then \mathcal{PKE} is IND-CCA secure.

We give a proof for this theorem for a PKE adversary A making exactly one encryption oracle query ($q_e = 1$).

More formally, we show for any IND-CCA adversary A against PKE with $q_e=1$, there exist adversaries B and C such that:

$$\text{Adv}_{\mathcal{PKE}}^{\text{IND-CCA}}(A) \leq 2 \cdot \text{Adv}_{\mathcal{KEM}}^{\text{IND-CCA}}(B) + \text{Adv}_{\mathcal{DEM}}^{\text{IND-CCA}}(C).$$

Moreover, B and C run in the same time as A . If A makes q_d decryption oracle queries, then B makes at most q_d decapsulation queries and C makes at most q_d decryption queries.

The KEM/DEM Composition Paradigm

Proof (sketch):

The proof involves a sequence of games G_0, G_1, G_2, G_3 .

G_0 : A plays the normal IND-CCA game for PKE, with $q_e=1$; let $c = (c_o, c_1)$ denote the response to A 's single encryption oracle query on (m_o, m_1) ; let b be the hidden bit chosen by A 's challenger. By construction:

$$(c_o, K_o) \leftarrow \text{KEM.Encap}(pk) \text{ and } c_1 \leftarrow \text{DEM.Enc}(K_o, m_b).$$

PKE dec(\cdot) oracle: as per the PKE scheme, using private key sk .

G_1 : A plays the normal IND-CCA game for PKE, with $q_e=1$; the response (c_o, c_1) to A 's encryption oracle query on (m_o, m_1) is constructed using "real" output from an IND-CCA challenger in the KEM security game. By construction:

$$(c_o, K_o) \leftarrow \text{KEM.Encap}(pk) \text{ and } c_1 \leftarrow \text{DEM.Enc}(K_o, m_b).$$

PKE dec(\cdot) oracle: on input (c'_o, c'_1) : if $c'_o \neq c_o$ then we use a KEM decap(\cdot) oracle from KEM security game to get a key K (or \perp) and then output $\text{DEM.Dec}(K, c'_1)$; otherwise, (when $c'_o = c_o$), we decrypt c'_1 using K_o .

The KEM/DEM Composition Paradigm

Proof (sketch):

G₂: As G₁ but the response (c_o, c_1) to A's encryption oracle query on (m_o, m_1) is constructed using "random" output K_1 from an IND-CCA challenger in the KEM security game. By construction:

$$(c_o, K_o) \leftarrow \text{KEM.Encap}(pk) \text{ and } c_1 \leftarrow \text{DEM.Enc}(K_1, m_b) \text{ where } K_1 \leftarrow_{\$} \mathcal{K}.$$

PKE dec(·) oracle: on input (c'_o, c'_1) : if $c'_o \neq c_o$ then we use the KEM decap(·) oracle on c'_o to get a key K (or \perp) and then output $\text{DEM.Dec}(K, c'_1)$; otherwise (when $c'_o = c_o$), we output $\text{DEM.Dec}(K_1, c'_1)$.

G₃: As G₂ but we now run KEM.KGen to get (sk, pk) , obtain c_o by running $\text{KEM.Encap}(pk)$, and obtain c_1 from a call to the $\text{enc}(\cdot)$ oracle in IND-CCA security game for the DEM (with hidden bit b). By construction:

$$(c_o, K_o) \leftarrow \text{KEM.Encap}(pk) \text{ and } c_1 \leftarrow \text{DEM.Enc}(K_1, m_b) \text{ where } K_1 \leftarrow_{\$} \mathcal{K}.$$

PKE dec(·) oracle: on input (c'_o, c'_1) : if $c'_o \neq c_o$ then we use sk ; otherwise (when $c'_o = c_o$), we use a call to the $\text{dec}(\cdot)$ oracle in DEM security game.

The KEM/DEM Composition Paradigm

Proof (sketch):

The proof involves a sequence of games G_0, G_1, G_2, G_3 .

From G_0 to G_1 : this is mostly a syntactic change, where we change how decryption is done (introducing a KEM challenger and its oracles).

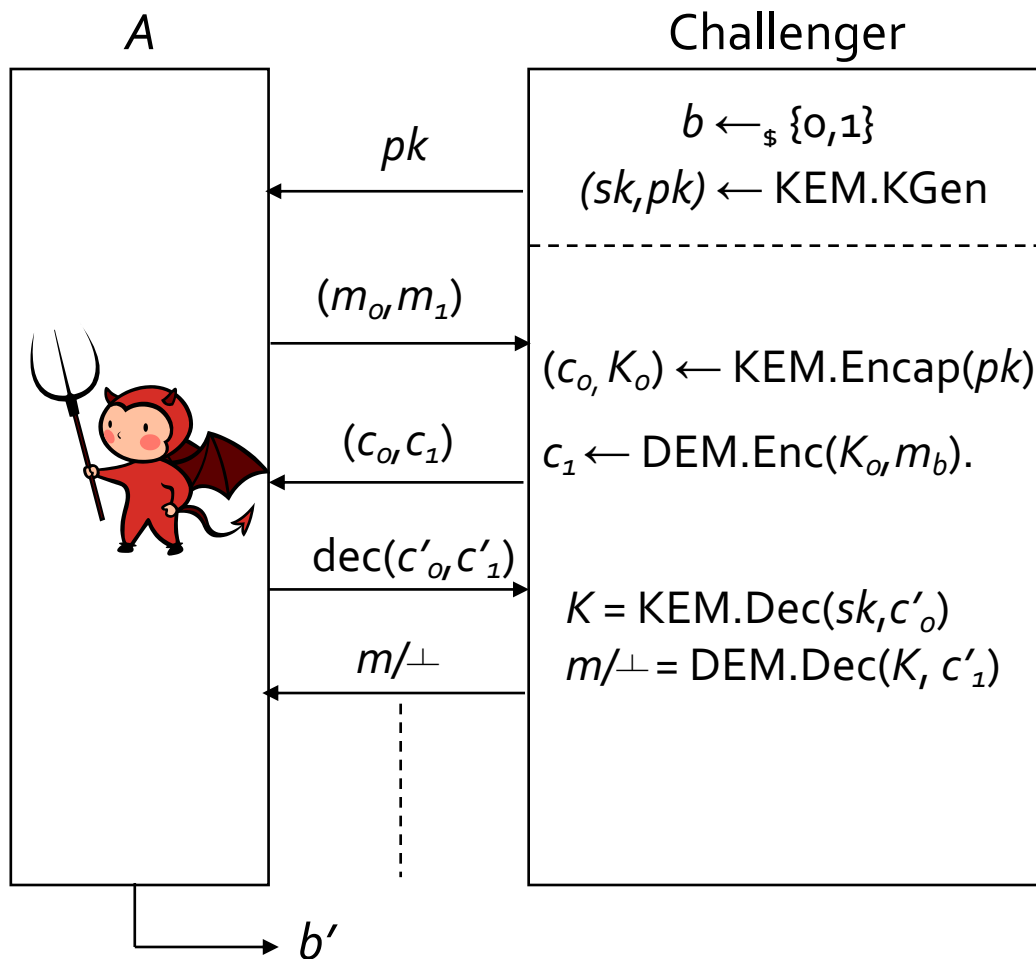
From G_1 to G_2 : we can construct an IND-CCA adversary B for the KEM that “bridges” between these games (G_1 corresponds to $b=0$ case, G_2 to $b=1$ case in KEM game). Doing this allows us to use K_0 in the KEM but **independent, random** K_1 in the DEM.

From G_2 to G_3 : this is again mostly a syntactic change, where we modify how encryption and decryption are done (using DEM oracles instead of KEM oracles).

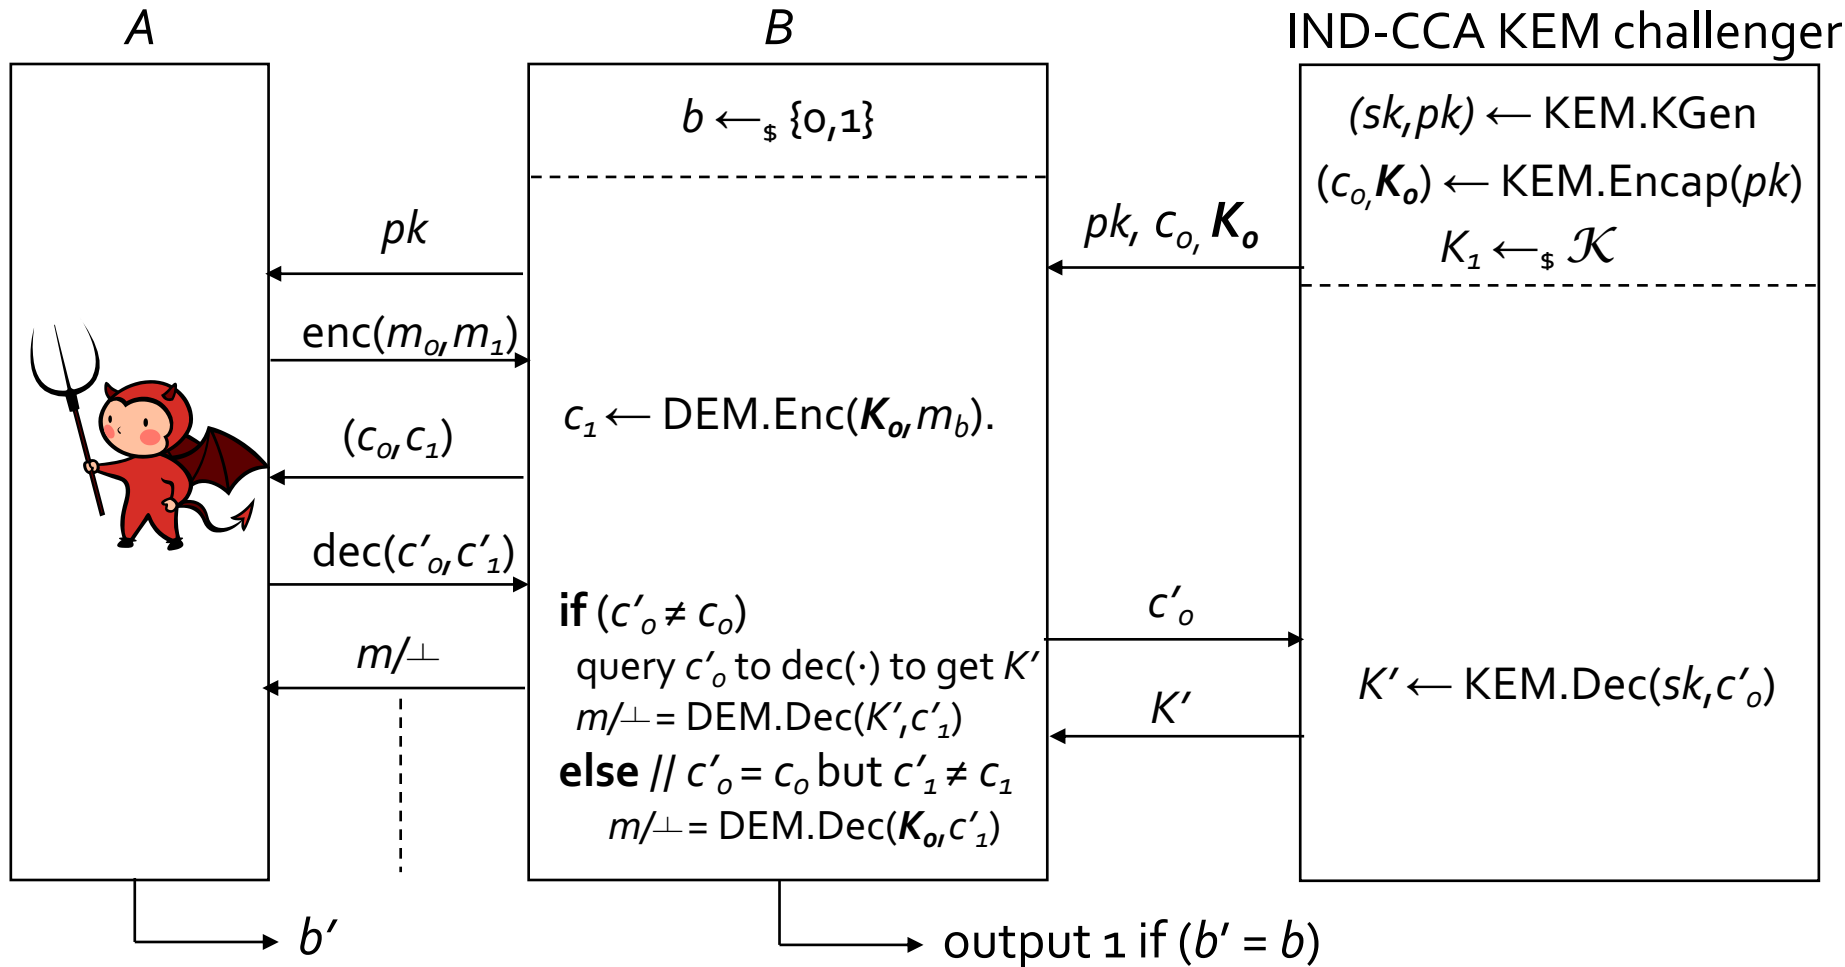
In G_3 : an adversary A who wins in this game can be used to construct an IND-CCA adversary C against the DEM.

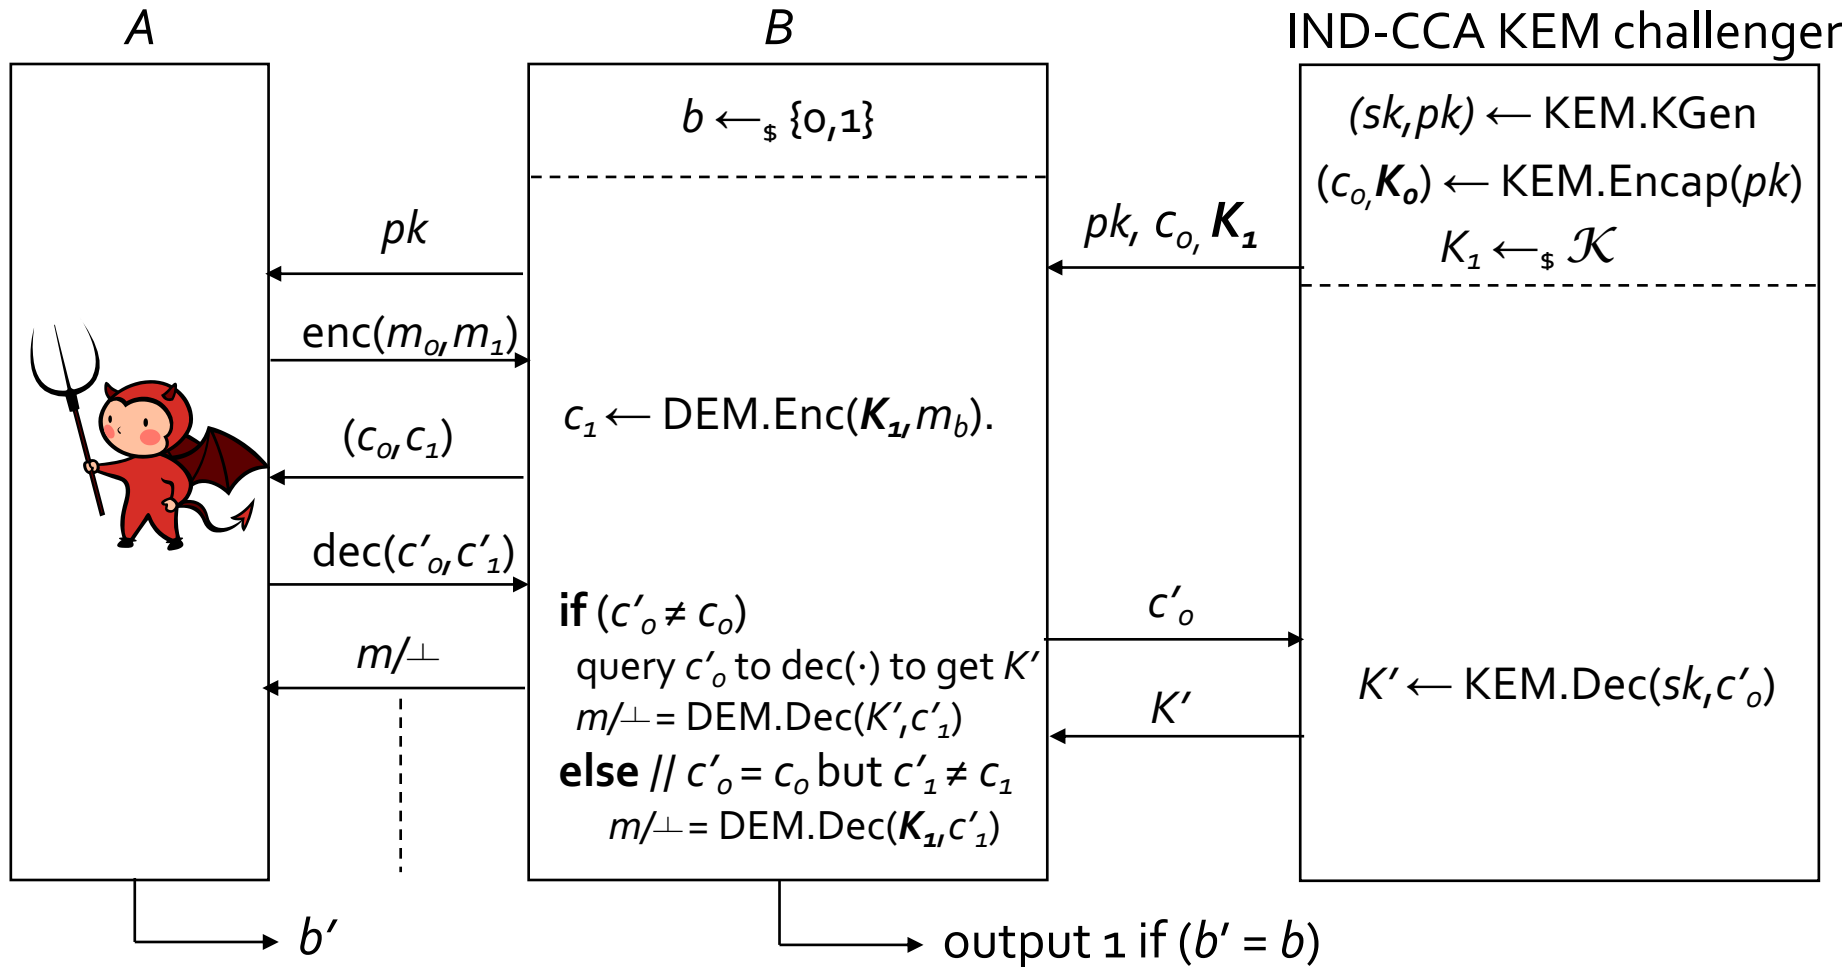
A standard computation involving probabilities of events W_i (the event that A outputs $b' = b$ in Game G_i) completes the proof.

G_o

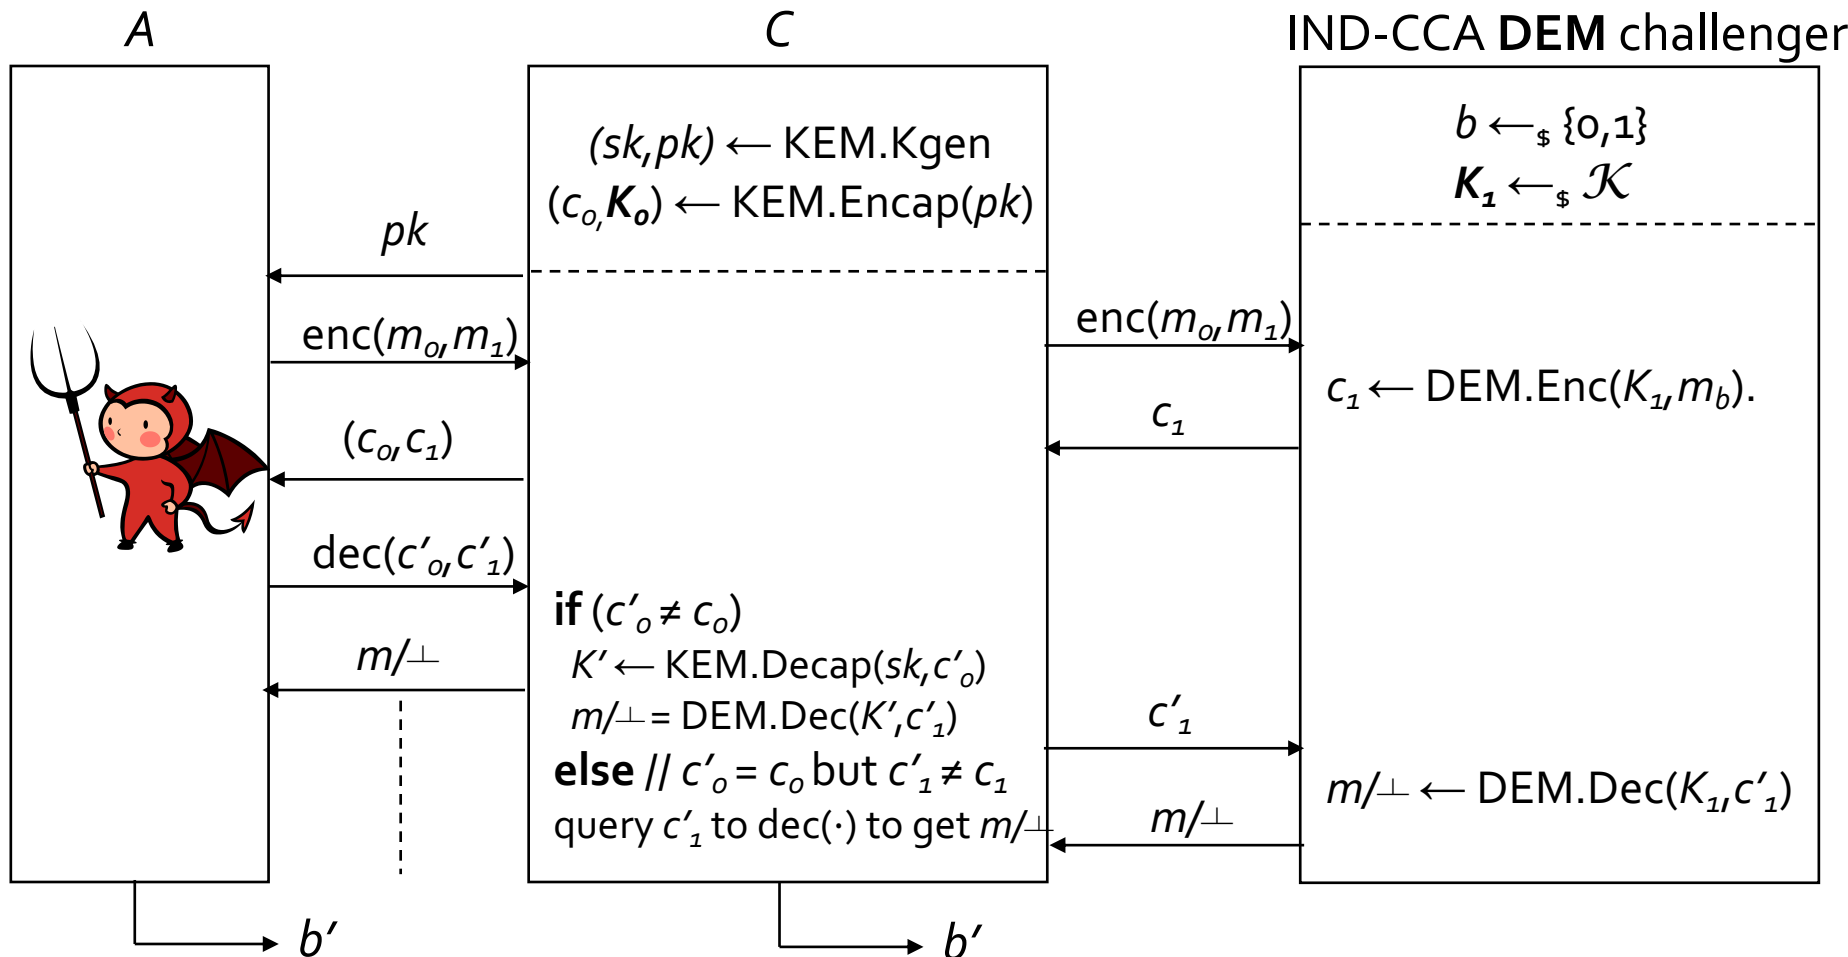


G_1



G_2 

G_3



The KEM/DEM Composition Paradigm

Notes:

- In game G_3 , C is an IND-CCA adversary against the DEM that only makes one encryption query (but many decryption queries).
 - This implies that only “one-time” security of the DEM is required.
- The proof is only for an IND-CCA adversary A against the PKE scheme that makes a single encryption query.
 - It can be extended to an adversary making q_e queries, at the cost of some complexity in the proof and some factors q_e in the bounds.
 - But one-time security of the DEM still suffices.
- We don’t have an integrity definition for PKE/KEMs; the public-key setting makes creating ciphertext/encapsulation forgeries trivial.

RSA Encryption

Recap: Textbook RSA

KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$.

Enc: inputs public key $pk = (e, N)$, plaintext $m \in [1, N-1]$; output $c = m^e \bmod N$.

Dec: inputs private key $sk = d$, ciphertext c ; output $m = c^d \bmod N$.

Recall:

- This scheme is not secure and must not be used in practice.
- It's not even randomised, so has no chance of even being IND-CPA secure.

Generating keys for RSA

KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$.

Generating large, random primes with $k/2$ bits:

- Needs a good source of randomness and an efficient primality test with low error rate.

Many things can go wrong, part 1:

“Mining your p 's and q 's”:

- $N_1 = p_1 q_1, N_2 = p_1 q_2 \Rightarrow p_1 = \gcd(N_1, N_2) \Rightarrow$ easy recovery of common prime p_1 and break of both keys.
- If M distinct RSA moduli are available, we can compute $O(M^2)$ pairwise gcds using an $O(M \log M)$ algorithm due to Bernstein.
- In 2012, this attack broke 0.50% of server public RSA keys on the Internet.
- Root cause: insufficient entropy in RSA key generation process.
- <https://factorable.net/weakkeys12.extended.pdf>.

Generating keys for RSA

Many things can go wrong, part 2:

ROCA attack on an over-optimised prime generation algorithm.

- Using p and q of special form to speed-up prime generation on smart-cards led to a major vulnerability in millions of smart-cards, including Estonian ID card system.
- https://crocs.fi.muni.cz/_media/public/papers/nemec_roca_ccs17_preprint.pdf

Take-aways:

- Implementations of even a 40-year-old algorithm can (and do) still get things wrong.
- Use a well-vetted library and standardised approaches to parameter selection.

Generating keys for RSA

KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$.

Solving $de = 1 \bmod (p-1)(q-1)$:

- Suppose e is coprime to $(p-1)(q-1)$.
- Running the extended Euclid algorithm yields integers s, t such that
$$e \cdot s + (p-1)(q-1) \cdot t = 1.$$
- Reduce modulo $(p-1)(q-1)$ to get:
$$e \cdot s = 1 \bmod (p-1)(q-1).$$
- Then take $d = s$.
- $e = 2^{16} + 1$ is often used in practice.
 - Faster encryption, no known weaknesses, e is prime and highly likely coprime to $(p-1)(q-1)$.
 - Many other tempting optimization lead to vulnerabilities (e.g. small d).

Using Chinese Remainder Theorem in RSA

Dec: inputs private key $sk = d$, ciphertext c ; output $m = c^d \bmod N$.

Textbook RSA decryption can be made faster by working mod p and mod q , and then combining the results using the Chinese Remainder Theorem (CRT).

Typically then include p , q and $u := q^{-1} \bmod p$ as part of the private key.

To decrypt:

- Set $f = d \bmod p-1$; $g = d \bmod q-1$
- Compute $m_p = c^f \bmod p$
- Compute $m_q = c^g \bmod q$
- Set $h = u(m_p - m_q) \bmod p$
- Output $m_q + hq \bmod N$

It is not too hard to show that this correctly implements RSA decryption.

Working with exponents f and g (mod p and mod q , respectively) makes the arithmetic operations much faster than direct implementation of $m = c^d \bmod N$.

Many things can go wrong, Part 3:

- In MEGA cloud storage system, $(d, p, q, u=q^{-1} \bmod p)$ are stored on server encrypted under a symmetric key K known only to the client **using ECB mode!**
- On login, client fetches ECB-encrypted (d, p, q, u) from server, decrypts to obtain (d, p, q, u) , recomputes $u'=q^{-1} \bmod p$ and tests if $u'=u$ to sanity check the key.
- But an adversarial cloud server can **overwrite** blocks of private key with chosen values, because the system exposes an encryption oracle for key K (due to a key reuse vulnerability).
- So the client tries to compute $(q')^{-1} \bmod p'$ where now p' and q' are controlled by the adversary.
 - This process succeeds if and only if q' is coprime to p' .
 - And if it fails an error message is returned to the server.
 - Based on presence/absence of the error message, information about coprimality of p', q' leaks to the adversary.
 - **Key trick:** make q' depend on selected blocks of d by exploiting malleability of ECB mode when doing private key overwriting.
 - Allows complete recovery of d using around $2^{11} - 2^{12}$ login attempts.
 - Full details at: <https://mega-caveat.github.io/>

Keysize requirements

- One way to break RSA is to recover the private key d from the public key (e, N) by first factorising N to find p, q and then solving $de = 1 \bmod (p-1)(q-1)$.
- Other attacks may be possible, but we have to *at least* make sure that factorising N is hard.
- The **Integer Factorisation Problem (IFP)** has been studied for thousands of years, and intensively since the 1970s (because of its importance in cryptography).
- The current best algorithm for solving IFP on a classical computer is called the **Number Field Sieve (NFS)**.
- This algorithm was invented in the early 1990s, and there has been no significant **algorithmic** improvement in IFP since then.
- But Moore's law + code improvements have had a significant effect.
- Running time + space required by NFS:

$$\exp[(c+o(1)) (\ln N)^{1/3} (\ln \ln N)^{2/3}] \quad \text{with} \quad c = (64/9)^{1/3}.$$

- Super-polynomial but sub-exponential in the bit-size of N .

Keysize requirements

- In 2015, a **512-bit RSA modulus** could be factored using NFS on Amazon EC2 in about 4 hours at a cost of USD75.
 - See <https://eprint.iacr.org/2015/1000>.
- Between 2009 and 2019, the biggest modulus publicly factored was a **768-bit RSA modulus**, one of the RSA challenges.
 - Using the equivalent of about 2,000 core years on 2.2 GHz CPUs.
 - See <https://eprint.iacr.org/2010/006.pdf>.
- Nov. 2019: RSA-240 (**795** bits), Boudot et al.: 900 core years – significantly less computation than the 2009 record!
- Feb. 2020: RSA-250 (**829** bits), Boudot et al.: 2700 core years on Intel Xeon Gold 6130 CPUs as a reference (2.1GHz), using CadoNFS software.
- See <https://lists.gforge.inria.fr/pipermail/cado-nfs-discuss/2020-February/001166.html> and <https://arxiv.org/abs/2006.06197>

RSA-250

2140324650240744961264423072839333563008614715144755017797
7549208814180234471401366433455190958046796109928518724709
1458768739626192155736304745477052080511905649310668769159
0019759405693457452230589325976697471681738069364894699871
578494975937497937

=

6413528947707158027879019017057738908482501474294344720811
6859632024532344630238623598752668347708737661925585694639
798853367

*

3337202759497815655622601060535511422794076034476755466678
4520987023841729210037080257448673296881877565718986258036
932062711

Keysize requirements

- It is conjectured that a **1024-bit modulus** would require about 2^{80} machine operations, hence providing 80 bits of security.
 - This is well beyond the capabilities of academic groups but probably within reach for a national security agency with a large budget.
- To achieve 128-bit security, the modulus N needs to be somewhere between **2048** and **3072** bits.
 - Estimates vary – BSI vs NIST vs ECRYPT.
 - See <https://www.keylength.com>.
 - Most websites still using RSA keys now use keys of 2048 bits, a few use 3072.
 - Question: Why not just “go large” on keysize?

Other attacks on RSA: Malleability

- We already know that textbook RSA is not **IND-CPA** secure, and we are going to need some kind of randomisation.
- But this is not all: textbook RSA also has a dangerous **malleability** property:
 - Suppose $c = m^e \bmod N$.
 - Choose s arbitrarily from $[1, N-1]$.
 - Consider $s^e \cdot c \bmod N$:
$$s^e \cdot c = s^e \cdot m^e = (s \cdot m)^e \bmod N.$$
 - So $s^e \cdot c \bmod N$ is an encryption of $s \cdot m \bmod N$ in textbook RSA!
 - This gives an active attacker the ability to modify an RSA plaintext in a controlled fashion, by carefully modifying ciphertexts.
 - This leads to various attacks against deployed systems.


Other attacks on RSA: Small e

- Using small exponents, such as $e = 3$, can speed up encryption.
- Suppose Alice has public key ($e=3$, N) and suppose Bob is only sending a very small message $m < N^{1/3}$.
- Then $c = m^3$ **over the integers**, i.e., no modular reduction has taken place.
- If Eve knows that m will be very small then she can recover the message m from the ciphertext c simply by taking cube roots over the integers (trivial, using Newton's method).
- Example:
 N has 1024 bits, m is an AES key with 128 bits, $e = 3$.
Then m^e has only $3 \times 128 = 384$ bits, much less than 1024 bits.
- To mitigate such attacks, **we need some kind of message padding**.
- Small d is also insecure even up to $d \approx N^{1/4}$ (Wiener's attack).

Padding for RSA

- In response to attacks like those we've discussed, special padding schemes for the RSA setting have been introduced.
- The three goals of such padding schemes are to:
 1. Introduce randomness into the message.
 2. Expand short messages to full size.
 3. Destroy algebraic relationships among messages to remove the malleability property.
- The ultimate aim should be to achieve IND-CCA security for RSA encryption.
- This has proved tricky to do, and even today the most widely-deployed RSA padding scheme, called PKCS#1 v1.5, does not achieve IND-CCA security.
- We do have padding schemes that enable us to achieve IND-CCA security, but they are not so widely deployed, e.g. RSA-OAEP.
- Why?

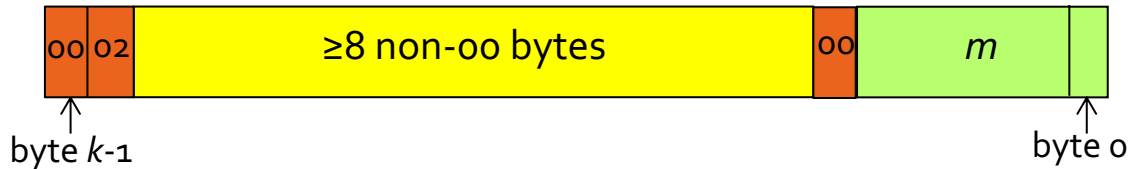
PKCS#1 v1.5 padding



- Byte-oriented encoding scheme for RSA.
- Assume N has k bytes; maximum message size is then $k - 11$ bytes.
- Message m placed in **least** significant bytes.
- Set first two most significant bytes to "00 02".
- Follow with **at least** 8 non-00 **random** bytes, then a 00 byte.
- Then apply the RSA operation to this encoded version of m :

$$c = \text{pad}(m)^e \bmod N.$$

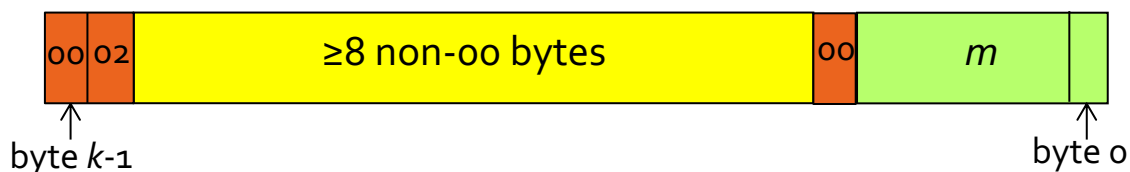
PKCS#1 v1.5 padding



Dec_d(c):


- Compute $m' = c^d \bmod N$.
- Check that m' begins with 00 02, reject if not.
- Check that m' then has at least 8 non-00 bytes, reject if not.
- Check that m' then has a 00 byte; reject if least significant byte is reached without finding one.
- Return as m all the bytes to the right of the 00 byte.
- Security?

Bleichenbacher's attack on PKCS#1 v1.5




- If an attacker sends a random string of bytes for decryption, then it has probability about 2^{-16} of having a valid padding format:
 - the first two bytes must be "00 02" – prob. 2^{-16} .
 - there are then at least 8 non-00 bytes – prob. $(255/256)^8 = 0.97$.
 - there is at least one "00" byte before the LSB is reached – prob. depends on key-length, higher for longer keys.
- Suppose now that the attacker has an **oracle** telling it whether the decryption algorithm succeeds or fails on an input c' .
- As with padding oracles for CBC mode, such an oracle is common in implementations.

Bleichenbacher's attack on PKCS#1 v1.5



- Attacker wants to decrypt $c = (00\ 02 || r || 00 || m)^e \bmod N$.
- Let's write $pad(m) = (00\ 02 || r || 00 || m)$ to denote the **padded** message.
- Attacker asks for decryption of $s^e \cdot c = (s \cdot pad(m))^e \bmod N$, for some choice of s .
- Most choices of s will lead to a decryption failure.
- With probability roughly 2^{-16} , decryption **succeeds**, and the adversary learns that $s \cdot pad(m) \bmod N$ begins with bytes "00 02".
- This gives a range for $s \cdot pad(m) \bmod N$ over the integers:
$$s \cdot pad(m) \bmod N \in [2B, 3B-1] \text{ where } B = 2^{8(k-2)}$$
- By gathering many such inequalities for different and carefully chosen values of s in an *adaptive* attack, attacker can recover $pad(m)$, and hence m .

Bleichenbacher's attack on PKCS#1 v1.5



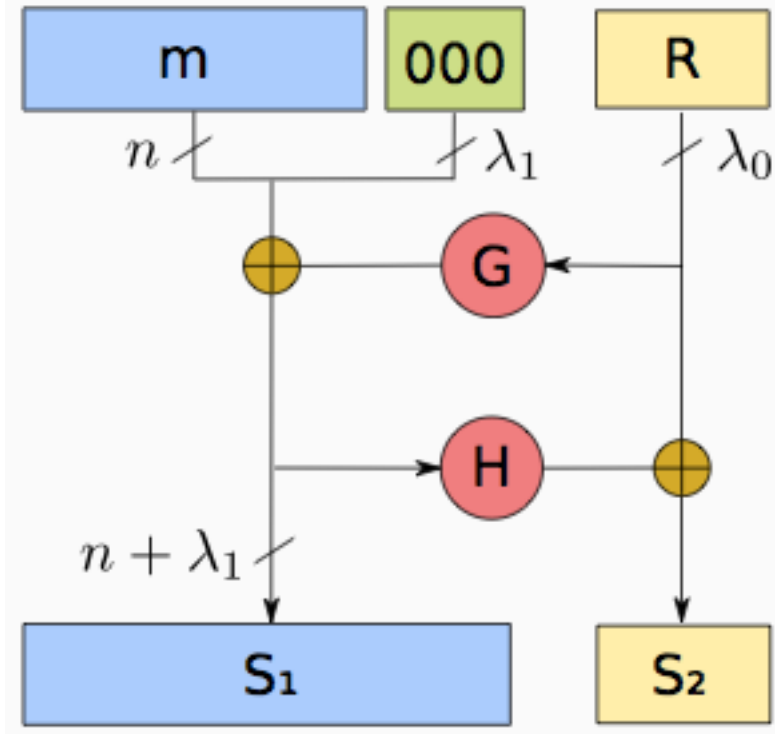
- Bleichenbacher's attack requires many decryption attempts, around 2^{20} in its original version.
- The attack has been improved a lot; many variants have been discovered.
- Typically 5-10k queries are needed for a 1024-bit modulus N in modern versions.
- Many real-world systems, including SSL/TLS, have been found to be vulnerable to the style of attack.
- For example, DROWN and ROBOT attacks on SSL/TLS.
 - <https://drownattack.com/>
 - <https://robotattack.org/>
- RSA-PKCS#1 v1.5 should be avoided but is still widely used in practice.
- A similar attack shows that RSA-PKCS#1 v1.5 is not IND-CCA secure.

RSA-OAEP

- OAEP is an alternative encoding scheme for RSA due to Bellare and Rogaway. OAEP stands for “Optimal Asymmetric Encryption Padding”.
- A version of it is widely standardized in PKCS#1 v2.1.
- **Setup of OAEP.** Suppose we are using λ -bit RSA moduli (e.g. $\lambda = 2048$).
- Let λ_o and λ_1 be chosen so that no adversary can perform 2^λ operations for $\lambda = \lambda_o$ or $\lambda = \lambda_1$, e.g. $\lambda_o = \lambda_1 = 128$.
- Set $n = \lambda - \lambda_o - \lambda_1$.
- Messages in RSA-OAEP are assumed to be n -bit strings.
- Let G be a hash function mapping λ_o bit strings to $n + \lambda_1$ bit strings.
- Let H be a hash function from $n + \lambda_1$ bit strings to λ_o bit strings.
- These hash functions should be **collision resistant**.


RSA-OAEP

- Let m be an n -bit string.
- Choose a random λ_0 -bit string R .
- Set $S_1 = (m \parallel 0^{\lambda_1}) \oplus G(R)$.
- Set $S_2 = R \oplus H(S_1)$.
- Form the bit-string $S = S_1 \parallel S_2$.
- Interpret S as a λ -bit integer and compute
$$c = S^e \bmod N.$$
- Decryption reverses these steps and checks that the result has λ_1 0-bits in the right position.
- Decryption fails if the 0-bits are not present; outputs m if they are.



Intuition for RSA-OAEP design

- RSA messages are randomised and have full length.
- Since S_1 and S_2 are both outputs of hash functions, the bit-string S looks random and algebraic relationships among the message components are broken up..
- Suppose you are given a decryption oracle. Because of the way the hash functions are combined, it's hard to come up with a ciphertext that decrypts to produce λ_0 zero bits in the correct position, without having done an encryption in the first place; so the decryption oracle essentially becomes useless.



- RSA-OAEP can be proven to be **IND-CCA** secure, though with strong assumptions on G and H , and under a strong number theoretic assumption (much stronger than assuming that factoring is hard). It's not perfect, but:

Use RSA-OAEP!

RSA.KEM: Building a simple KEM from RSA

RSA.KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$. Let $H: \{0, \dots, N-1\} \rightarrow \{0, 1\}^k$ be a hash function.

RSA.Encap: inputs public key $pk = (e, N)$; $s \leftarrow_{\$} \{0, \dots, N-1\}$; return (c, K) where

$$c = s^e \bmod N; K = H(s)$$

RSA.Decap: inputs private key $sk = d$ and encapsulation c .

$$\text{compute } s = c^d \bmod N;$$

$$\text{output } K = H(s).$$

Theorem:

RSA.KEM is IND-CCA secure in the **Random Oracle Model** provided the **RSA inversion problem** is hard.

The RSA Inversion Problem

1. Challenger C runs RSA KeyGen to produce pair (sk, pk) with
 $sk = d, pk = (e, N)$.
2. C selects $x \leftarrow_{\$} \{0, \dots, N-1\}$; sets $y = x^e \bmod N$.
3. A is given input (N, e, y) by C ; A runs and outputs some value x' .
 A wins if $x' = x$.

- Strictly speaking, this version of the RSA Inversion problem is relative to the choice of KeyGen that we use.
- If A can factor N then it can solve the RSA inversion problem.
- The reverse implication is open, but no algorithm faster than factoring N is known for solving RSA inversion in general.
- For analysis of RSA.KEM, see extra slides.
- RSA.KEM is a much better scheme than RSA-OAEP, but is not widely used in practice.

Introducing the Discrete Logarithm Setting

The discrete logarithm setting

- So far, all of our public key schemes have been in the RSA setting:
 - Public/verification keys involve integers N that are the product of two primes, p, q , along with an integer e .
 - Private/signing keys are integers d such that $de = 1 \bmod (p-1)(q-1)$.
 - Security depends on hardness of factoring and related problems (RSA inversion).
- The discrete logarithm setting provides another platform for carrying out public key cryptography.
 - After introducing this setting, we'll look at Diffie-Hellman key exchange, and finally public key encryption.

The discrete logarithm setting

- Assume p and q are large primes and q divides $p-1$.
- So we can write $p = kq + 1$ for some integer k ; so $k = (p-1)/q$.
- (So p and q do **not** have the same flavour as in RSA!).
- Often, but not always, $k=2$.

Toy example: $p = 29$, $q = 7$, $k = 4$.

- In reality, for 128-bit security, p will need to have 3072 bits and q will need to have (at least) 256 bits.
- Typical real-world deployments: p has 1024 or 2048 bits, q has (at least) 160 bits.

The discrete logarithm setting

- Now we pick h , a random integer mod p , and compute $g = h^{(p-1)/q} \bmod p (= h^k \bmod p)$.
 - If $g = 1 \bmod p$, we try again (we will succeed with very high probability)
- **Fact 1:** if $g \neq 1 \bmod p$, then the q powers of g , namely $G_q = \{g, g^2, g^3, \dots, g^q\}$ are all distinct mod p .
- **Fact 2:** $g^q = 1 \bmod p$.
- **Fact 3:** if we multiply together two elements in the set G_q , we obtain a third element that is also in the set G_q .
- In combination, these facts mean that the set G_q forms a **group** of order q ; the group operation is multiplication mod p .
 - G_q is a cyclic group (everything is a power of g).
 - We say that g is a generator of G_q .
 - The number of elements in G_q is q , a prime.

The discrete logarithm setting

Toy example (ctd):

- $p = 29, q = 7, k = 4$ ($27 = 4 \cdot 7 + 1$).
- $h = 2; g = h^{(p-1)/q} = 2^4 = 16 \bmod 29$.
- $G_q = \{g, g^2, g^3, \dots, g^q\}$:
 - So $g^2 = 16^2 = 256 = 24 \bmod 29$; and:
 - $g^3 = 16^3 = 16^2 \times 16 = 24 \times 16 = 384 = 7 \bmod 29$;
 - $g^4 = 16^4 = 16^3 \times 16 = 7 \times 16 = 112 = 25 \bmod 29$;
 - $g^5 = 16^5 = 16^4 \times 16 = 25 \times 16 = 400 = 23 \bmod 29$;
 - $g^6 = 16^6 = 16^5 \times 16 = 23 \times 16 = 368 = 20 \bmod 29$;
 - $g^7 = 16^7 = 16^6 \times 16 = 20 \times 16 = 320 = 1 \bmod 29$;
- Hence $G_q = \{16, 24, 7, 25, 23, 20, 1\}$.
- And, for example, $24 \times 7 = g^2 \times g^3 = g^5 = 23 \bmod 29$.

The discrete logarithm problem (DLP)

- The set $G_q = \{g, g^2, g^3, \dots, g^q\}$ of powers of $g \bmod p$ is a **cyclic group of prime order q with generator g** .
- It forms a subgroup of the integers mod p under multiplication.
- We can write $g^q = 1 = g^0$ and so take $G_q = \{1=g^0, g^1, g^2, g^3, \dots, g^{q-1}\}$.

The discrete logarithm problem in G_q :

Let (p, q, g) be as above. Set $y = g^x \bmod p$, where x is a uniformly random value in $\{0, 1, \dots, q-1\}$. **Find x .**

(Think of value x as being the logarithm of $y = g^x \bmod p$ when the base of logarithms is g .)

Solving the DLP

- The DLP has received intense analysis from mathematicians and computer scientists for nearly 40 years.
 - Several different algorithms exist, but they are beyond the scope of this course.
 - This is an active area of research, with recent breakthroughs in related settings.
- The **Function Field Sieve** has complexity of the form:
$$\exp[(1+o(1)).(32/9)^{1/3} (\log p)^{1/3}(\log \log p)^{2/3}]$$
which is sub-exponential (but super-polynomial) in $\log p$.
- The **Pollard- p algorithm** and its variants has complexity of the form: $O(q^{1/2})$, which is exponential in $\log q$.
- To achieve 80-bits of security against both algorithms, we need q to have 160 bits and p to have 1024 bits.
- To achieve 128-bits of security against both algorithms, we need q to have 256 bits and p to have about 3072 bits.

Cryptography in the discrete logarithm setting

- Basic message so far: we can choose p and q big enough so as to make the DLP in G_q sufficiently hard.
- But we have to keep an eye on developments in algorithms for solving the DLP.
- Similar to RSA setting: what if someone comes up with a better factoring algorithm? Or a large quantum computer?
- Question now is: can we use the DLP to build cryptosystems (in particular, encryption schemes)?
- We will look at Diffie-Hellman, then ElGamal.

Diffie-Hellman Key Exchange

Diffie-Hellman Key Exchange (DHKE)

- DHKE is a public key method for agreeing on a shared secret (which can be used as a session key, for example).
- Introduced in famous 1976 paper by Diffie and Hellman that launched public key cryptography.
 - <https://www-ee.stanford.edu/~hellman/publications/24.pdf>
- We will use the Discrete Logarithm setting.
- Recall: $p = k.q + 1$, and g generates G_q , a cyclic group of prime order q in the set of integers modulo p :

$$G_q = \{g^0 = 1, g^1, g^2, \dots, g^{q-1}\}.$$

Diffie-Hellman Key Exchange

- In Diffie and Hellman's original presentation, we have a collection of n users.
- All users make use of the same set of public parameters (p, q, g) .
- User U_i picks x_i uniformly at random from $\{0, 1, \dots, q-1\}$.
- User U_i 's private key is x_i ; its public key is $Y_i = g^{x_i} \bmod p$.
- All the public keys Y_1, \dots, Y_n are assumed to be available in a public, authentic directory.
 - In the 1970s, telephone directories were common!

Diffie-Hellman Key Exchange

To compute a shared value between user U_i and user U_j :

- User U_i computes $K_{ij} = (Y_j)^{x_i} = g^{x_j \cdot x_i} = g^{x_i \cdot x_j} \bmod p$;
- User U_j computes $K_{ji} = (Y_i)^{x_j} = g^{x_i \cdot x_j} \bmod p$.
- Key point: the values K_{ij} and K_{ji} are the same.
- We should not use K_{ij} directly as a key since it is large integer mod p .
- Instead, we should derive key(s) from it using a Key Derivation Function (KDF).
- For now, think of KDF as just being a hash function.

Diffie-Hellman Key Exchange

- Notice that no interaction is needed between any pair of users in order to establish a key.
- Instead, the users are assumed to have access to the public, authentic database of **public** keys Y_1, \dots, Y_n .
- We have **non-interactive key exchange (NIKE)**.
- **Question 1:** How is NIKE different from PKE?
- **Question 2:** Can we get from NIKE to PKE?

Security of Diffie-Hellman Key Exchange

- An attacker can see all public keys Y_1, \dots, Y_n and would like to compute some (or all) of the shared values K_{ij} (or the keys derived from them).
- The underlying algorithmic problem can be stated as follows:

Given (p, q, g) , and the values $g^a \bmod p$, $g^b \bmod p$, for uniformly random a, b , find $g^{ab} \bmod p$.
- This is the Computational Diffie-Hellman Problem (CDHP).

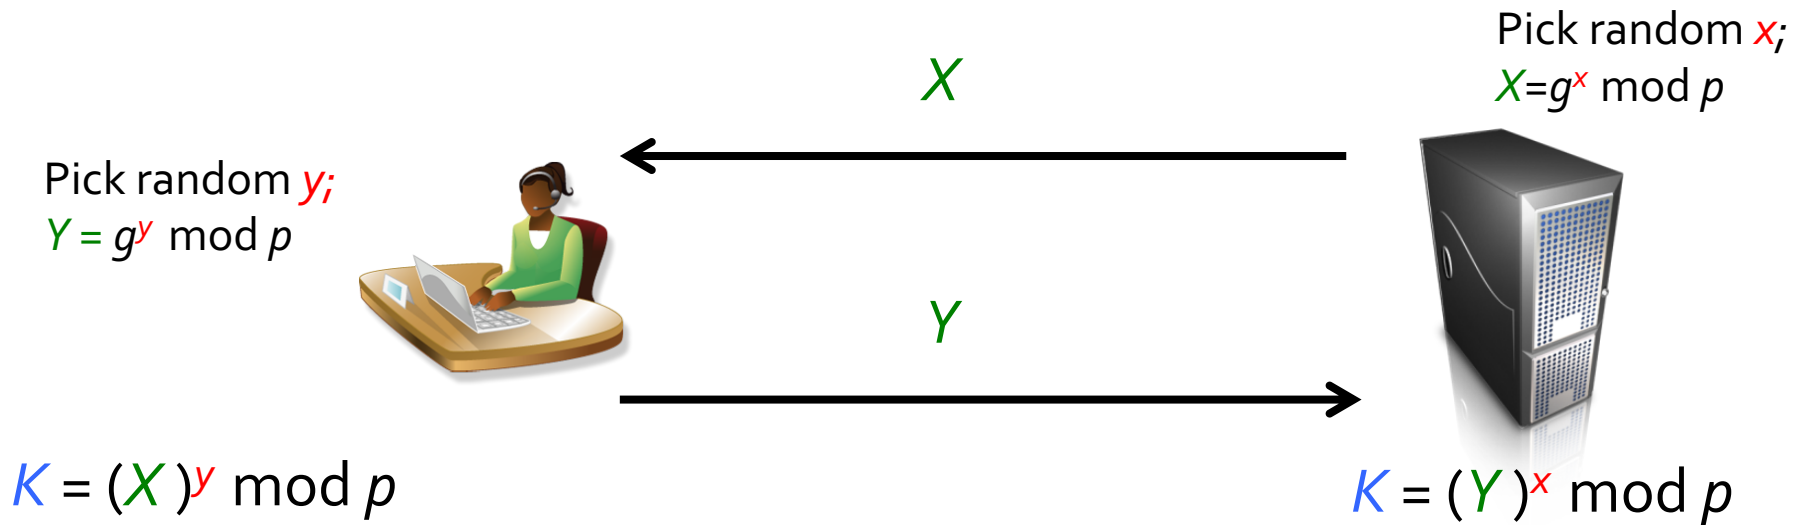
Security of Diffie-Hellman Key Exchange

- The Computational Diffie-Hellman Problem (CDHP).

Given (p, q, g) , and the values $g^a \bmod p$, $g^b \bmod p$, for uniformly random a, b , find $g^{ab} \bmod p$.

- If we can efficiently solve the DLP, then we can efficiently solve CDHP: find a from $g^a \bmod p$ by solving DLP, then compute $(g^b)^a \bmod p$.
- So CDHP is not harder than DLP, and could be easier.
- We don't know in general whether CDHP is equivalent to DLP, but this is widely believed to be so, and known to be so in special cases (den Boer 1988, Maurer-Wolf 1999).
- This helps us choose secure parameters for Diffie-Hellman: just assume CDHP is as hard as DLP.

A modern view of Diffie-Hellman Key Exchange



A modern view of Diffie-Hellman Key Exchange

- We simplify to just two parties, Alice and Bob.
- Alice and Bob first agree on parameters g , p and q , select fresh random private values x , y , then exchange the corresponding values X , Y over a public communications channel.
- We regard the private values x , y and the public values X , Y as being *ephemeral*.
 - That is, they are generated once and used only once, setting up a fresh shared value $K = g^{xy} \bmod p$ each time.
- It now seems that any two parties can securely agree upon a key without having had any previous association!

MITM Attack on Diffie-Hellman Key Exchange

But an active man-in-the-middle attacker can modify the public values in transit.

- For example, the attacker can change X to a value $g^z \bmod p$ for which he knows z .
- The attacker can then compute the key that Alice would compute on receipt of $g^z \bmod p$, namely, $(g^z)^y = (g^y)^z = (Y)^z \bmod p$.
- And similarly in the other direction.
- So an active attacker can completely compromise the security of “modern” Diffie-Hellman key exchange.
- Was that the case for the NIKE version?

Preventing MITM Attack on Diffie-Hellman Key Exchange

- So the modern version requires the **authenticity** of the public values X, Y to be protected in transit.
- How can this be done?
- Two possible solutions: use a MAC or use digital signatures.

Q: If we had a key for a MAC, why do we need to agree a new key via Diffie-Hellman?

A: Forward security: even if the MAC key was later revealed, the new key from Diffie-Hellman would still be secure (more in later lectures).

Q: What additional requirements does the use of signatures bring?

A: We now need an assurance of the authenticity of the signature verification keys. This is typically done using certificates, CAs, and PKI.

From Diffie-Hellman to ElGamal Encryption

From Diffie-Hellman Key Exchange to Public Key Encryption

- Diffie-Hellman paper appeared in 1976.
- Introduced the *concept* of public key encryption, but contained no concrete example of a public key encryption scheme.
- Included Diffie-Hellman key exchange (with a public database).
- So can we get from DHKE to PKE?
- Yes we can!
- We next define the ElGamal public key encryption scheme, dating from 1985.

ElGamal Public Encryption Scheme

Public parameters (p, q, g) as usual (can be shared amongst many users).

KeyGen: pick x uniformly at random from $\{0, 1, \dots, q-1\}$. Set public key to be $X = g^x \bmod p$ and private key to be x .

Enc: given public key X and message M (assumed to be encoded as an element of G_q):

1. Pick r uniformly at random from $\{0, 1, \dots, q-1\}$. Set $Y = g^r \bmod p$.
2. Compute $Z = X^r \bmod p$.
3. Output the ciphertext $C = (Y, M \cdot Z \bmod p)$.

ElGamal Public Encryption Scheme

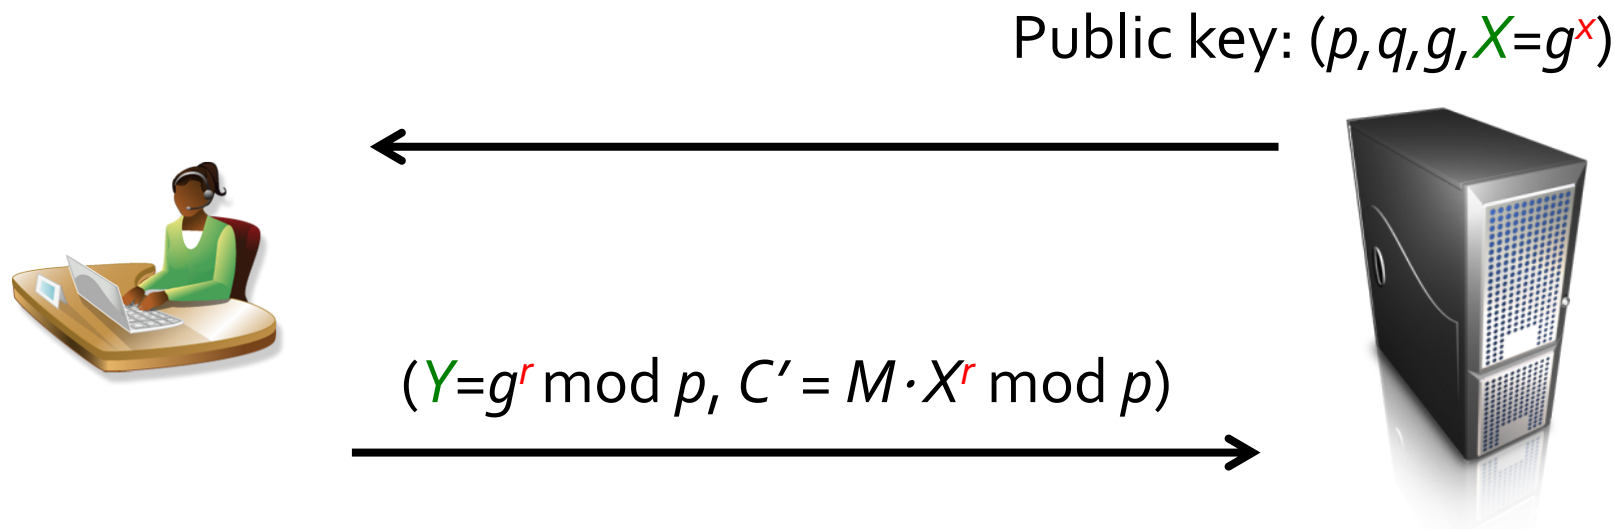
Dec: given private key x and ciphertext $C = (Y, C')$:

1. Check that Y is in G_q , return “fail” if not.
2. Compute $Z' = Y^x \bmod p$.
3. Output $M = C' \cdot (Z')^{-1} \bmod p$.

Correctness of the scheme:

$$\begin{aligned} C' \cdot (Z')^{-1} &= M \cdot Z \cdot (Z')^{-1} = M \cdot X^r \cdot (Y^x)^{-1} \bmod p \\ &= M \cdot (g^x)^r \cdot ((g^r)^x)^{-1} \bmod p \\ &= M \cdot g^{xr} \cdot g^{-xr} \bmod p \\ &= M \end{aligned}$$

Relationship Between Diffie-Hellman and ElGamal



- We effectively have a Diffie-Hellman key exchange, involving long-term key pair (x, X) and one-time key pair (r, Y) , with the sender choosing (r, Y) and including the public value Y as part of the ciphertext.
- We use the shared Diffie-Hellman value $X^r = g^{xr} \bmod p$ as an encryption mask.
- The encrypting party need not have a “registered” public key to use the scheme (as should be the case for PKE!).

Security of ElGamal Public Encryption Scheme

IND-CPA security for ElGamal encryption can be proven based on the hardness of a variant of the CDH Problem, called the Decisional Diffie-Hellman Problem:

Given (p, q, g) , and uniformly random values a, b, c from $\{0, 1, \dots, q-1\}$, distinguish the triple (g^a, g^b, g^{ab}) from the triple (g^a, g^b, g^c) .

Informally:

- Ciphertext includes $M \cdot Z \bmod p$ where $Z = g^{xr} \bmod p$.
- If (g^x, g^r, g^{xr}) is indistinguishable from (g^x, g^r, g^c) , then we can replace $Z = g^{xr}$ by $Z = g^c$ in the construction of ciphertext.
- But, since g^c is uniformly random in G_q , it follows that $M \cdot Z \bmod p$ is uniformly random in G_q as well.
- Hence M is perfectly hidden!

Formally: do a game hopping proof with G_0 the normal IND-CPA game and G_1 the game in which the challenge ciphertext uses g^c in place of g^{xr} .

DHIES

DHIES

ElGamal PKE has two significant disadvantages:

- It is not **IND-CCA** secure (find an attack exploiting the homomorphic nature of ElGamal encryption!)
- It is inconvenient to work with messages that have to be encoded as elements of G_q .
- DHIES (Diffie-Hellman Integrated Encryption Scheme) addresses both of these problems in a PKE scheme.
- One can also define a KEM based on ElGamal and use the KEM/DEM paradigm.

DHIES

Public parameters (p, q, g) as usual, H a hash function with suitable output domain.

KeyGen: pick x uniformly at random from $\{0, 1, \dots, q-1\}$. Set public key to be $X = g^x \bmod p$ and private key to be x .

Enc: given public key X and message M (a bit-string):

1. Pick r uniformly at random from $\{0, 1, \dots, q-1\}$. Set $Y = g^r \bmod p$.
2. Compute $Z = X^r \bmod p$.
3. Set $K = H(Z, X, Y)$.
4. Split K into K_e and K_m .
5. Compute $C' = \text{EtM}(M)$ using keys K_e and K_m for encryption and MAC, respectively.
6. Output the ciphertext $C = (Y, C')$.

DHIES

Dec: given private key x and ciphertext $C = (Y, C')$:

1. Check that Y is in G_q , return “fail” if not.
2. Compute $Z = Y^x \bmod p$.
3. Set $K = H(Z, x, Y)$.
4. Split K into K_e and K_m .
5. Decrypt C' using keys K_e and K_m for encryption and MAC, respectively.
6. Output “fail” if step 5 fails, otherwise output the message returned in step 5.

Correctness again relies on: $x^r = Z = Y^x \bmod p$, that is, that we implicitly have a Diffie-Hellman key exchange.

DHIES

- DHIES is **IND-CCA** secure in the Random Oracle Model.
- The security proof relies on the hardness of a variant of the Computational Diffie-Hellman Problem (assuming we use an IND-CPA symmetric encryption scheme and a SUF-CMA MAC scheme).
- It's really an instance of the KEM/DEM paradigm: we are sending information Y that allows a shared key K to be computed, and then using that key K in a symmetric encryption scheme.
- We could replace “**EtM**” with any **AE** scheme.
- Can go nonce-based, use AEAD, etc.

From Diffie-Hellman Key Exchange to Public Key Encryption

- An elliptic-curve variant of DHIES also exists, called ECIES.
 - It has smaller ciphertext sizes, because of the more compact representation of group elements (g^x) in the elliptic curve setting.
 - Security relies on the hardness of the Elliptic Curve Diffie-Hellman (ECDH) Problem.
- These schemes are standardised in ANSI X9.63, IEEE P1363a, ...
- ECIES will be used in the next generation of EMV “chip and PIN” card standards.
- We’ll return to ECIES when we look at elliptic curve cryptography.

Homework

- Read Chapters 10-12 of Boneh-Shoup for many more details, constructions and proofs.
- Work on exercise sheets 8, 9.
- Prepare for this week's lab.

Appendix: Some Basic Number Theory

Basic Number Theory: Remainders, factors,...


- We use the notation \mathbf{N} to denote the non-negative integers:

$$\mathbf{N} = \{0, 1, 2, \dots\}.$$

- We write $a \bmod b$ for the “remainder of a on division by b ”, i.e. the integer r in $\{0, \dots, b-1\}$ such that $a = q \cdot b + r$ for some q .
- For example,
 - $9 \bmod 7 = 2$, because $9 = 1 \cdot 7 + 2$;
 - $21 \bmod 7 = 0$, because $21 = 3 \cdot 7 + 0$.
- We say “ a is a factor of b ” if a divides b with remainder 0.
- For example,
 - 3 is a factor of 6, but
 - 4 is not a factor of 6.

Basic Number Theory: Modular arithmetic

- Modular arithmetic refers to arithmetic done with the numbers in $\{0, \dots, b-1\}$ (for some b), where we always reduce our results to remainders on division by b .
- For example, with $b = 7$, we have:
 - $6 + 4 = 10 = 3 \bmod 7$.
 - $6 \cdot 4 = 24 = 7 \cdot 3 + 3 = 3 \bmod 7$.
 - $6^4 = 6 \cdot 6 \cdot 6 \cdot 6 = (6^2)^2 = 36^2 = 1^2 = 1 \bmod 7$.
- Important principle: we never need to work with numbers much bigger than b if we are doing arithmetic mod b : we can just reduce mod b as we go along.
- Note that if b is composite, then we can multiply two non-zero numbers mod b and get zero!
 - For example, with $b = 12$ we have $6 \cdot 4 = 24 = 2 \cdot 12 + 0 = 0 \bmod 12$.
- This is different from normal arithmetic....



Basic Number Theory: Primes and gcds

- A number greater than 1 with only itself and 1 as factors is called a **prime number**. For example, 13 is prime, but 21 is composite (= not prime).
- The **greatest common divisor** (gcd) of two numbers is the largest number that is a factor of both numbers. For example, $\text{gcd}(4, 6) = 2$.
- Two numbers are said to be **relatively prime** if their gcd is 1.

Basic Number Theory: Modular inverses

- Over the rational numbers we define the (multiplicative) inverse of a , denoted a^{-1} as the number such that $a \cdot a^{-1} = 1$.
- Modulo some number p , we can also define (modular) inverses.
 - For example, $5 \cdot 5 \bmod 6 = 25 \bmod 6 = 1 \bmod 6$.
 - Hence $5^{-1} = 5 \bmod 6$.
- All numbers $0 < a < p$ such that $\gcd(a, p) = 1$ have an inverse modulo p .
- But if $0 < a < p$ is such that $\gcd(a, p) > 1$, then a will not have an inverse modulo p .
- If p is a prime, then $\gcd(a, p) = 1$ for all $0 < a < p$ (why?), and so every number a in $\{1, \dots, p-1\}$ will have an inverse mod p .
- The numbers modulo a prime p form a field, which means that addition, multiplication and division are defined and work for all inputs “as expected”.

Basic Number Theory

- As an example, let's make the addition and multiplication tables mod 5.

+	0	1	2	3	4
0	0	1	2	3	4
1	1	2	3	4	0
2	2	3	4	0	1
3	3	4	0	1	2
4	4	0	1	2	3

·	0	1	2	3	4
0	0	0	0	0	0
1	0	1	2	3	4
2	0	2	4	1	3
3	0	3	1	4	2
4	0	4	3	2	1

- Notice the pattern along rows and columns of addition table.
- Notice how every element has a multiplicative inverse except 0.
- 1 and 4 are self-inverse.
- As an exercise, make the same tables for modulus 6.

Basic Number Theory: The Euclidean algorithm

- Given $a, b \in \mathbf{N}$ we want to compute $\gcd(a, b)$.
- Euclid's algorithm is based on recursive application of the formulae:
$$\gcd(a, b) = \gcd(b, a \bmod b) \quad (1)$$
$$\gcd(a, 0) = a \quad (2)$$
- Notice that $0 \leq a \bmod b < b$, so the numbers decrease in size at every step.
- Example:
$$\begin{aligned} & \gcd(21, 14) \\ &= \gcd(14, 21 \bmod 14) \quad (\text{using (1)}) \\ &= \gcd(14, 7) \\ &= \gcd(7, 14 \bmod 7) \quad (\text{using (1)}) \\ &= \gcd(7, 0) \\ &= 7. \quad (\text{using (2)}) \end{aligned}$$

Basic Number Theory: The extended Euclidean algorithm

- The **extended Euclidean algorithm** outputs integers r, s, t where $r = \gcd(a, b)$ and $s \cdot a + t \cdot b = r$.
- In other words, it gives an expression for $\gcd(a, b)$ as a linear combination of a and b over the integers.
- It works like the Euclidean algorithm but keeps track of additional information as we go along.
- It computes a series of coefficients r_i, s_i, t_i such that:

$$r_i = s_i \cdot a + t_i \cdot b.$$

- In each step we compute

$$r_{i+1} = r_{i-1} - q_i \cdot r_i \quad \text{where} \quad 0 \leq r_{i+1} < |r_i|$$

$$s_{i+1} = s_{i-1} - q_i \cdot s_i$$

$$t_{i+1} = t_{i-1} - q_i \cdot t_i$$

- We initialise the algorithm with $r_0 = a, r_1 = b$ and $s_0 = 1, s_1 = 0, t_0 = 0, t_1 = 1$.
- Eventually $r_k = \gcd(a, b)$ and $\gcd(a, b) = s_k \cdot a + t_k \cdot b$.

Basic Number Theory: The extended Euclidean algorithm

- The **extended Euclidean algorithm** outputs integers r, s, t where $r = \gcd(a, b)$ and $s \cdot a + t \cdot b = r$.
- Example:

i	r_i	q_i	s_i	t_i	$r_i = s_i \cdot a + t_i \cdot b$
0	$a = 14$	—	1	0	$1 \cdot 14 + 0 \cdot 21 = 14$
1	$b = 21$	—	0	1	$0 \cdot 14 + 1 \cdot 21 = 21$
2	14	0	1	0	$1 \cdot 14 + 0 \cdot 21 = 14$
3	7	1	-1	1	$-1 \cdot 14 + 1 \cdot 21 = 7$
4	0	2	3	-2	$3 \cdot 14 - 2 \cdot 21 = 0$

Using the extended Euclidean algorithm

- The numbers s and t are often useful for applications.
- For example, running the extended Euclid's algorithm on a with $0 < a < p$ and a prime p gives s and t such that:

$$a \cdot s + p \cdot t = 1.$$

- Reducing this equation mod p , we get:

$$a \cdot s = 1 \bmod p.$$

- So Euclid's algorithm allows us to find the inverse of $a \bmod p$.
- Recall that in RSA key generation, we need to a pair (e, d) such that:
$$e \cdot d = 1 \bmod (p-1)(q-1).$$
- This can also be done using the extended Euclidean algorithm:
 - Choose e at random in the range $[1, (p-1)(q-1)]$.
 - With high probability, $\gcd(e, (p-1)(q-1)) = 1$ and then running Euclid's algorithm gives us the required d .

Correctness of textbook RSA

Let $N = pq$ and suppose $ed = 1 \bmod (p-1)(q-1)$.

Hence $(p-1)(q-1) \mid ed - 1$.

So write $ed = 1 + k(p-1)(q-1)$.

Let m be coprime to N . Then m is coprime to p and q .

By Fermat's little theorem, $m^{p-1} = 1 \bmod p$.

Then $m^{ed} = m^{1 + k(p-1)(q-1)} = m \cdot (m^{p-1})^{k(q-1)} = m \cdot 1^{k(q-1)} = m \bmod p$.

Hence p divides $m^{ed} - m$.

By symmetry q also divides $m^{ed} - m$. Then $N=pq$ also divides $m^{ed} - m$.

Hence $m^{ed} = m \bmod N$.

This completes the correctness proof.

(What about the case where m is not coprime to N ?)

Appendix: Building a Simple KEM from RSA

Building a simple KEM from RSA

RSA.KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$. Let $H: \{0, \dots, N-1\} \rightarrow \{0, 1\}^k$ be a hash function.

RSA.Encap: inputs public key $pk = (e, N)$; $s \leftarrow_{\$} \{0, \dots, N-1\}$; return (c, K) where

$$c = s^e \bmod N; K = H(s)$$

RSA.Decap: inputs private key $sk = d$ and encapsulation c .

$$\text{compute } s = c^d \bmod N;$$

$$\text{output } K = H(s).$$

Theorem:

The above scheme RSA.KEM is IND-CCA secure in the **Random Oracle Model** provided the **RSA inversion problem** is hard.

The RSA Inversion Problem

1. Challenger C runs RSA KeyGen to produce pair (sk, pk) with

$$sk = d, pk = (e, N).$$

1. C selects $x \leftarrow_{\$} \{0, \dots, N-1\}$; sets $y = x^e \bmod N$.
2. A is given input (N, e, y) by C ; A runs and outputs some value x' .
3. A wins if $x' = x$.

- NB $x = y^d = y^{1/e} \bmod N$, so x is the e -th root of $y \bmod N$.
- Strictly speaking, this version of the RSA Inversion problem is relative to the choice of KeyGen that we use.
- If A can factor N then it can solve the RSA inversion problem.
- The reverse implication is open, but no algorithm faster than factoring N is known for solving RSA inversion in general.

The Random Oracle Model (ROM)

- The ROM is a strong abstraction of hash functions.
- Formalized by Bellare and Rogaway (1993).
- In short, in the ROM, we model a hash function H as a *random* function (from a given domain to a given range).
- An adversary A cannot then compute H for itself but must instead “outsource” all its computations of H to its challenger in the form of oracle queries.
- This gives a security reduction in a proof the power to “inspect” queries made by A , possibly extracting useful information from them.
- It also allows the reduction to “program” specific values into its H -oracle replies.
- We can also argue that any output $H(x)$ is uniformly random unless and until A queries H on input x .
- We will see this in action in our analysis of RSA.KEM ahead.

The Random Oracle Model (ROM)

- Of course, real hash functions like SHA-256 are fixed functions and not random oracles.
- So we must **instantiate** the random oracle somehow in order to implement any scheme whose analysis use the ROM.
- **Heuristic step:** replace the random oracle with a fixed hash function.
- There are many arguments for and against using the ROM in security proofs.
- Pro: it's a useful tool, especially in the public key setting.
- Con: it's not formally sound.
- Pro: none of the examples showing this are very natural.
- Pro: it enables more efficient schemes (in comparison to the *standard* model).
- A lot of research effort has gone into closing the gap between standard-model-secure schemes and ROM-secure schemes.

Building a simple KEM from RSA (recap)

RSA.KGen: generates a pair of random primes p, q of some bit-size $k/2$, and integers d, e such that $de = 1 \bmod (p-1)(q-1)$. Set $N = pq$. Output key pair (sk, pk) where $sk = d, pk = (e, N)$. Let $H: \{0, \dots, N-1\} \rightarrow \{0, 1\}^k$ be a hash function.

RSA.Encap: inputs public key $pk = (e, N)$; $s \leftarrow_{\$} \{0, \dots, N-1\}$; return (c, K) where

$$c = s^e \bmod N; K = H(s)$$

RSA.Decap: inputs private key $sk = d$ and encapsulation c .

compute $s = c^d \bmod N$;

output $K = H(s)$.

An RSA inversion instance!

Uniformly random unless adversary queries H on s .

Proof Intuition for simple KEM from RSA – 1

Claim:

RSA.KEM is IND-CCA secure in the ROM assuming the hardness of RSA inversion.

Proof – intuition for IND-CPA security

We show that from any IND-CPA adversary A against the KEM, we can build an adversary B solving the RSA inversion problem.

Given an RSA inversion challenge (N, e, y) , B sets $pk = (N, e)$ and uses y to construct challenge encapsulation c^* in the KEM security game for A .

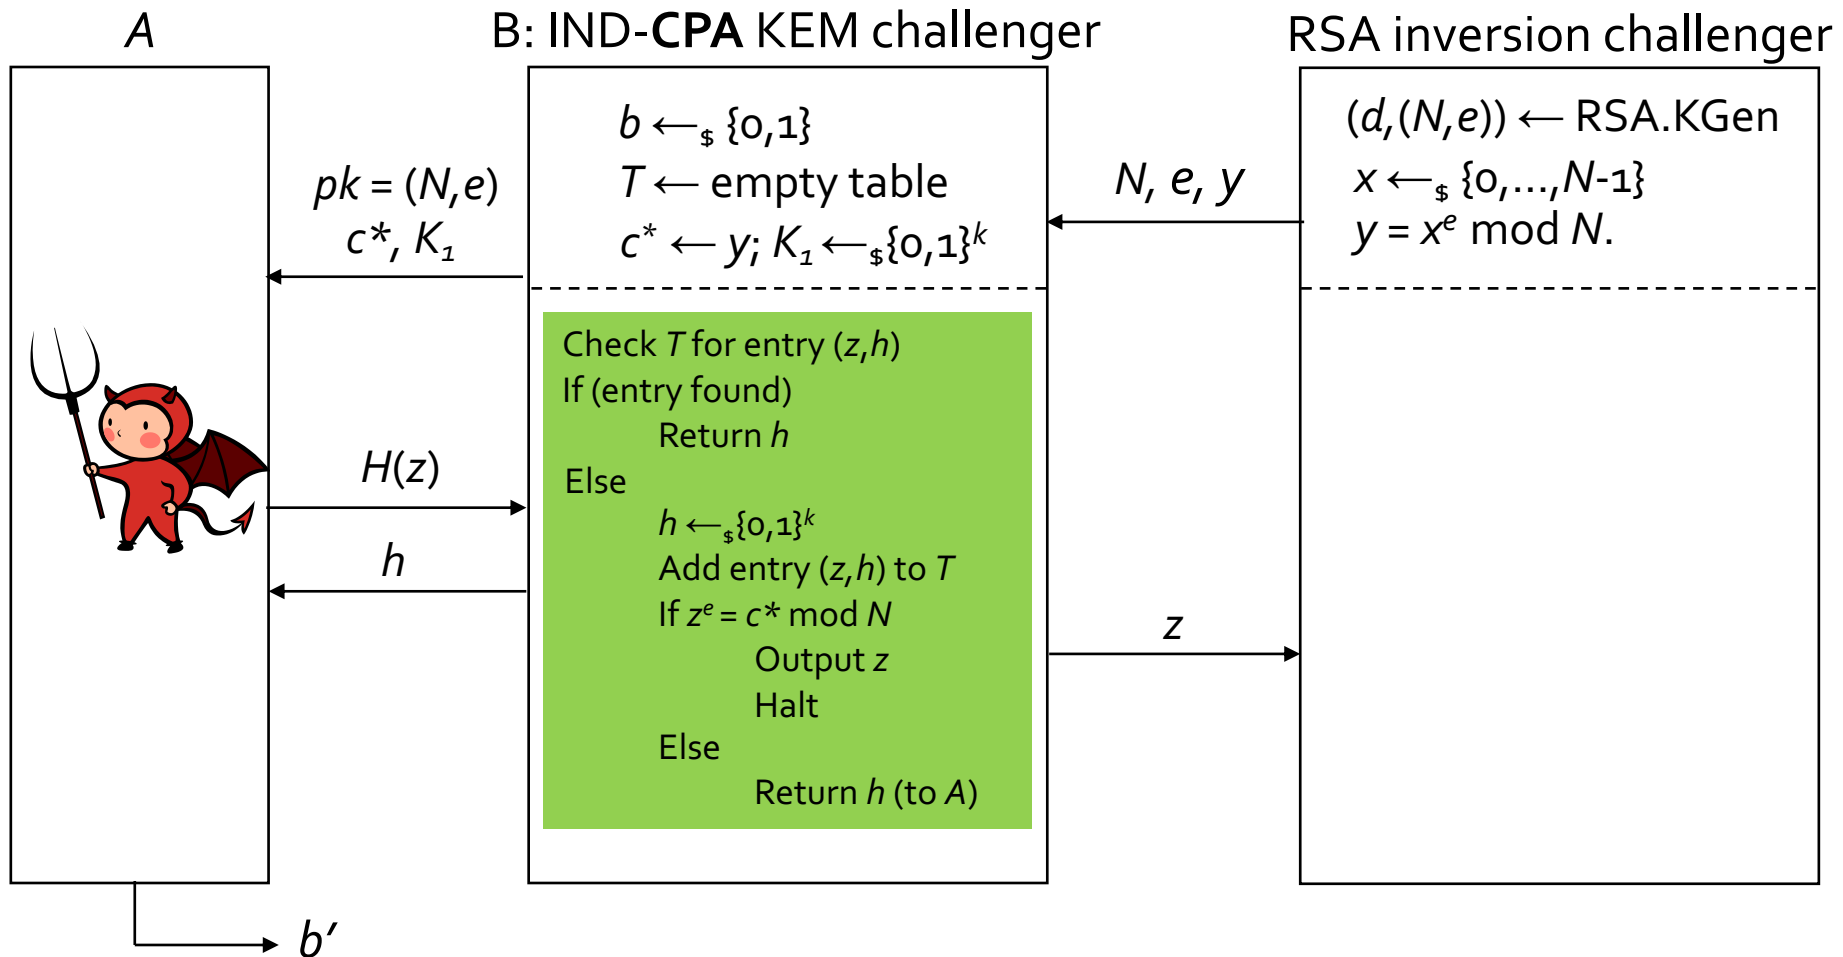
- Key K_0 then equals $H((c^*)^d \bmod N) = H(y^d \bmod N) = H(y^{1/e} \bmod N) = H(x)$.
- Key K_1 is set to random by B .

Recall that A 's challenger (played by B) gives (pk, c^*, K_b) to A ; to win in the KEM security game, A must be able to distinguish K_0 from K_1 , i.e. distinguish $H(x)$ from random.

A can't do this unless it queries H on x at some point, since otherwise $H(x)$ is uniformly random and has an identical distribution to K_1 .

If A does query H on x , then B can find x amongst the queries made by A ; in fact, A can check each H -query z made by A and test if $z^e = y \bmod N$.

Sketch of reduction for IND-CPA security



Proof Intuition for simple KEM from RSA – 2

- What about IND-CCA security?
- The crucial point is that B must provide a consistent simulation of $\text{dec}(\cdot)$ and $H(\cdot)$ queries for A , at least until x is queried (after that, we don't care – the simulation can go wrong).
- B “patches the random oracle” using a table of triples of one of two forms: (z, c, h) or $(?, c, h)$.
- **On $\text{dec}(c)$ query:**
 - Check for a table entry (z, c, h) ; if one exists, respond with h .
 - Otherwise, choose random h and add entry $(?, c, h)$ to the table; this implicitly defines $H(z) = h$ where $z = c^{1/e} \bmod N$ is unknown (hence “?” in place of z in the table).
- **On H query with input z :**
 - Check if there is already an entry (z, c, h) in the table; if so, return h (basic consistency).
 - Check if $z^e = c \bmod N$ for some entry $(?, c, h)$ in the table.
 - If yes, then complete table entry to (z, c, h) ; return h (maintain consistency with dec queries).
 - If not, select a random h , set $c = z^e \bmod N$ and add a new entry (z, c, h) ; return h .
- Working out all the details is a bit messy, so we omit them.
- B 's advantage is tightly related to that of A but its running time is higher (due to cost of maintaining table consistency).

Incomplete sketch of reduction

